

# Class 7: Intro to Machine Learning

Xinyu Wen (A17115443)

## Table of contents

K-means clustering . . . . .	4
Hierarchical Clustering . . . . .	8
<b>1. PCA of UK food data</b>	<b>10</b>
Hands on with Principal Component Analysis (PCA) . . . . .	10
Data import . . . . .	10
PCA to the rescue! . . . . .	14
Plot PC1 vs PC2 . . . . .	15
<b>2. PCA of RNA-seq data</b>	<b>20</b>
ggplot . . . . .	24
Optional: Gene loadings . . . . .	27


Today we will explore unsupervised machine learning methods including clustering and dimensionality reduction methods.

Let's start by making up some data (where we know there are clear groups) that we can use to test out different clustering methods.

We can use the `rnorm()` for

```
hist(rnorm(n=3000, mean = 3))
```

### Histogram of rnorm(n = 3000, mean = 3)



Make data `z` with two “clusters”

```
rnorm(30, mean=-3)
```

```
[1] -3.7433153 -4.3217792 -1.7935740 -4.0083938 -2.1423993 -4.8050819  
[7] -3.1227694 -3.8793063 -1.6400976 -2.4738805 -2.5137105 -3.2710059  
[13] -4.3814283 -2.7573337 -3.8191915 -3.4313084 -3.9037321 -3.7413905  
[19] -3.2917719 -2.1486427 -3.1573811 -3.2697133 -2.6017514 -4.7756139  
[25] -1.5591695 -0.9643865 -2.1290850 -2.4372159 -2.9797937 -2.4223979
```


```
rnorm(30, mean= +3)
```

```
[1] 1.9656432 2.6165892 2.4576659 4.0839193 2.5460514 2.2198801  
[7] 2.6294766 3.9199817 2.4131573 1.7513112 3.7873296 1.8843090  
[13] 3.7294562 2.7017269 3.1007470 4.8384023 3.7593797 -0.1232246  
[19] 4.1205845 6.2365091 2.0469441 3.7537896 3.0279897 3.3793664  
[25] 4.2383501 2.7450149 1.1914912 2.2843404 4.2715554 2.1194014
```

```
x <- c(rnorm(30, mean=-3),  
rnorm(30, mean= +3))
```

```
hist(x)
```


**Histogram of x**



```
x <- c(rnorm(30, mean=-3),  
rnorm(30, mean= +3))  
  
z <- cbind(x=x, y=rev(x))  
head(z)
```

```
      x      y  
[1,] -2.9295538 3.460984  
[2,] -0.8919533 4.394508  
[3,] -3.8755243 2.447696  
[4,] -2.6230031 4.457529  
[5,]  0.4548172 2.731386  
[6,] -2.0539523 3.883956
```

```
plot(z)
```



How big is z

```
nrow(z)
```

```
[1] 60
```

```
ncol(z)
```

```
[1] 2
```

## K-means clustering

The main function in “base” R for K-means clustering is called `kmeans()`.

```
k <- kmeans(z, centers = 2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

x	y
---	---

```
1 -2.696343  2.983828
2  2.983828 -2.696343
```

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 72.73951 72.73951
(between_SS / total_SS = 86.9 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
attributes(k)
```

\$names

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

\$class

```
[1] "kmeans"
```

Q. How many points lie in each cluster? (aka. size of cluster?)

k\$size

[1] 30 30

Q. What component of our results tells us about the cluster membership (i.e. which point likes in which cluster)?

```
k$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. Center of each cluster?

```
k$centers
```


	x	y
1	-2.696343	2.983828
2	2.983828	-2.696343

Q. Put this result info together and make a little “base R” plot of our clustering result. Also add the cluster center points to this plot.

```
plot(z, col = "blue")
```




```
plot(z, col=c(1, 2))
```



Plot colored by cluster membership:


```
plot(z, col=k$cluster)
points(k$centers, col="blue", pch = 15)
```



Q. Run kmeans on our input z and define 4 clusters making the same result visualization plot as above (plot of z colored by cluster membership).

```
k4 <- kmeans(z, center = 4)

plot(z, col=k4$cluster)
points(k$centers, col="purple", pch = 15)
```



## Hierarchical Clustering


The main function in base R for this is called `hclust()`. It will take as input a distance matrix (key point is that you can't just give your "row" data as input. You have to first calculate a distance matrix from your data).

```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:  
`hclust(d = d)`




```
plot(hc)
abline(h=10, col = "red")
```



```
grps <- cutree(hc, h=10)
```

```
plot(z, col = grps)
```



## 1. PCA of UK food data

### Hands on with Principal Component Analysis (PCA)

Let's examine a 17-dimensional data containing details of food consumption in the UK. Are these countries different? How?

#### Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1) #Q2, first way
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209

Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q1. How many rows and columns are in your new data frame named X? What R functions could you use to examine this?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

```
dim(x)
```

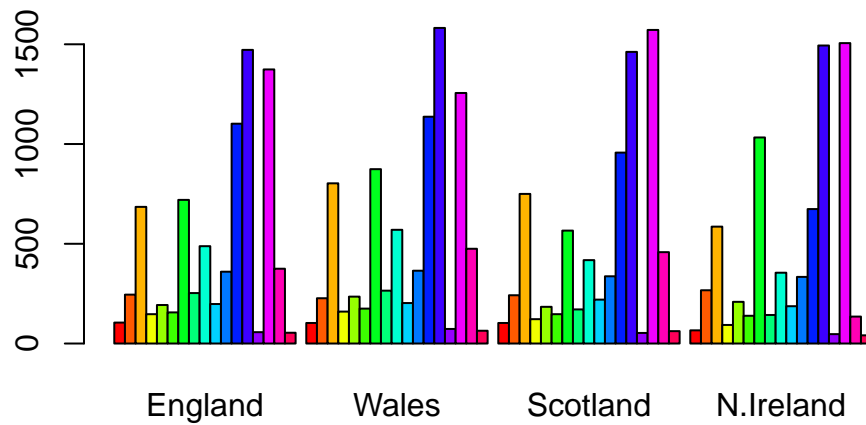
```
[1] 17 4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
#The second way
#rownames(x) <- x[,1]
#x <- x[,-1]
#head(x)
```


I like the first way better, because it eliminates first column and read the file in one line. However, if I do not know which column of the file I should eliminate, it would be better to print the original table first, then use the second way to eliminate the column.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

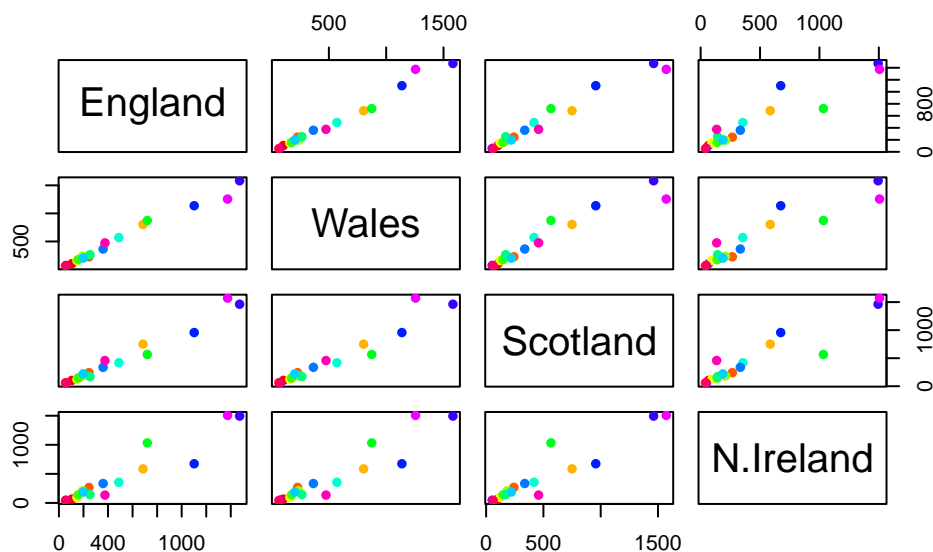
```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Set beside as `False`.

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(nrow(x)), pch =16)
```



If a point lies on the diagonal of a plot, it means the two countries have similar consumption of that food.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The main difference is that N. Ireland consumes much more fresh potatoes.

Looking at these types of “pairwise plots” can be helpful but it does not scale well. There must be a better way.....

## PCA to the rescue!

The main function for PCA in base R is called `prcomp()`. This function wants the transpose of our input data - ie the important foods in as columns and the countries as rows.

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is in our PCA result object `pca`.

```
head(pca$x)
```


	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

The `pca$x` result object is where we will focus first as this details how the countries are related to each other in terms of our new “axis” (aka. “PCs”, “eigenvectors”, etc.)


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

### Plot PC1 vs PC2

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"), pch =16, xlab="PC1",
```




```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "darkgreen"))
```





We can look at the so called PC “loadings” result object to see how the original foods contribute to our new PCs. (ie, how the original variables contribute to our new better PC variable)

```
pca$rotation[,1]
```

Cheese	Carcass_meat	Other_meat	Fish
-0.056955380	0.047927628	-0.258916658	-0.084414983
Fats_and_oils	Sugars	Fresh_potatoes	Fresh_Veg
-0.005193623	-0.037620983	0.401402060	-0.151849942
Other_Veg	Processed_potatoes	Processed_Veg	Fresh_fruit
-0.243593729	-0.026886233	-0.036488269	-0.632640898
Cereals	Beverages	Soft_drinks	Alcoholic_drinks
-0.047702858	-0.026187756	0.232244140	-0.463968168
Confectionery			
-0.029650201			

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```


```
## or the second row here...
```

```
z <- summary(pca)
```

```
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00


```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



```
## Lets focus on PC1 as it accounts for > 90% of variance
```


```
par(mar=c(10, 3, 0.35, 0))
```

```
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
#Loading plot for PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



Soft drinks and fresh potatoes are the two prominent food groups. PC2 tells us that the consumption of fresh potatoes and soft drinks have opposite trends. If fresh potatoes are larges consumed, soft drinks are not. And vise versa.

## 2. PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10. How many genes and samples are in this data set?

```
genes <- nrow(rna.data)
genes
```

```
[1] 100
```


```
samples <- ncol(rna.data)
samples
```

```
[1] 10
```

There are 100 genes and 10 samples.

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un-polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.345e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

```
plot(pca, main="Quick scree plot")
```




```
## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```


```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```



```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```




### ggplot

```
library(ggplot2)

df <- as.data.frame(pca$x)


# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```





```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

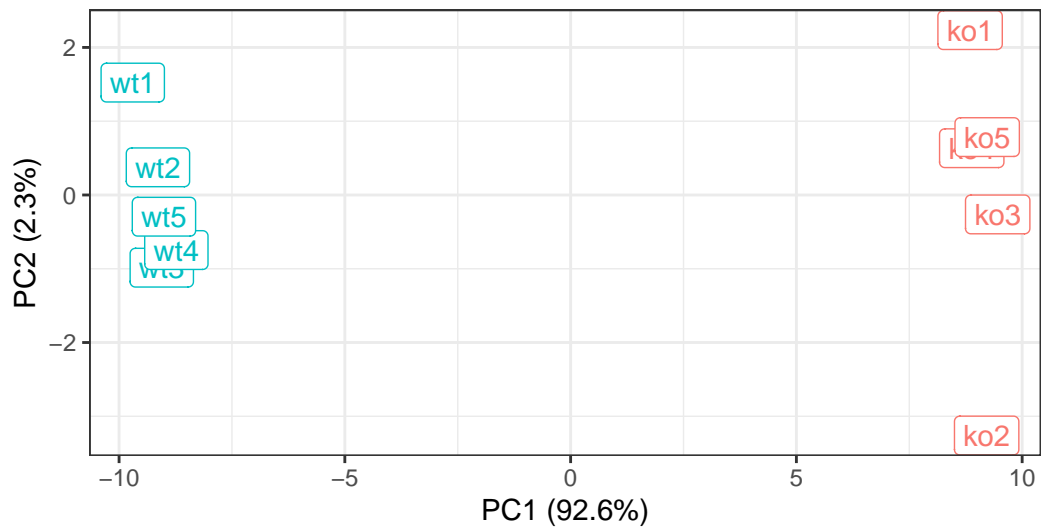
p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
  theme_bw()
```

## PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data

### Optional: Gene loadings

```
loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
[1] "gene100" "gene66" "gene45" "gene68" "gene98" "gene60" "gene21"
[8] "gene56" "gene10" "gene90"
```